SkillAssure

# Docker Storage

SkillAssure

# File System

- On installation, docker creates its folder structure at var/lib/docker
- This is where Docker stores all its data by default.

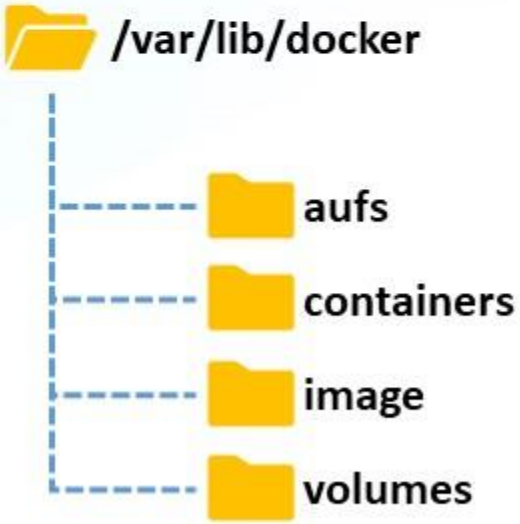

SkillAssure

# Layered Architecture

## Dockerfile

```
FROM Ubuntu

RUN apt-get update && apt-get -y install python

RUN pip install flask flask-mysql

COPY . /opt/source-code

ENTRYPOINT FLASK_APP=/opt/source-code/app.py  flask run
```

## Dockerfile2

```
FROM Ubuntu

RUN apt-get update && apt-get -y install python

RUN pip install flask flask-mysql

COPY app2.py /opt/source-code

ENTRYPOINT FLASK_APP=/opt/source-code/app2.py  flask run
```

## docker build Dockerfile –t prabhav/my-custom-app


## docker build Dockerfile2 –t prabhav/my-custom-app2

SkillAssure

# Layered Architecture

- All these layers are created when we build the Docker image.
- These are called Image Layers
- Once the build is complete, contents of these layers cannot be modified.
- These layers are Read Only
- These layers can be modified by initiating a new build

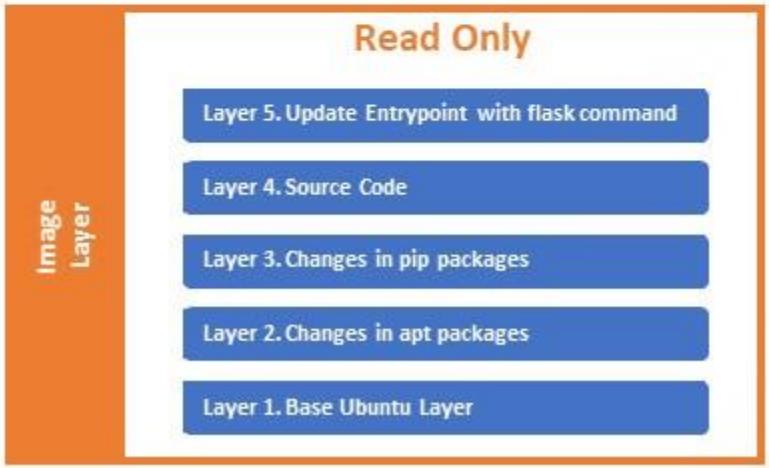

SkillAssure

# Layered Architecture

- When we run a container, docker creates a container from these layers and creates a new writable layer on top of the image layer.
- The writable layer is called as container layer
- The writable layer is used to store data created by the container
  - e.g. log files, temporary files or just any file modified by the user on that container.
- The life of this layer is only as long as the container is alive.
- When the container is destroyed this layer and all the changes stored in it are also destroyed.

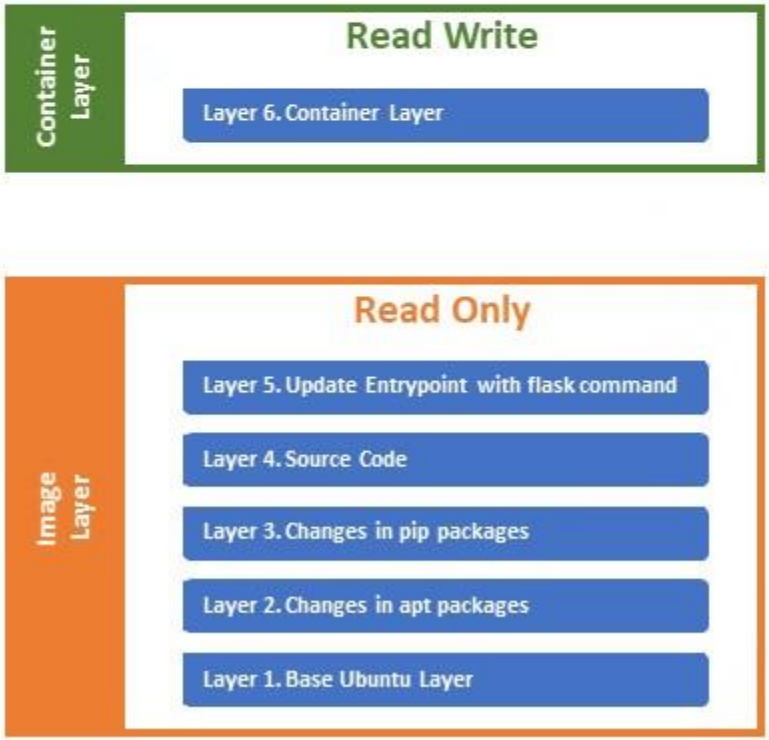

SkillAssure

# Layered Architecture

- Let's say we login to the newly created container and create a new file - temp.txt.
- It will get created in the container layer which is read and write.


- Our application code is in Image layer.
- If I wish to modify the application code
- How will the modification work? As our Image layer is read only.

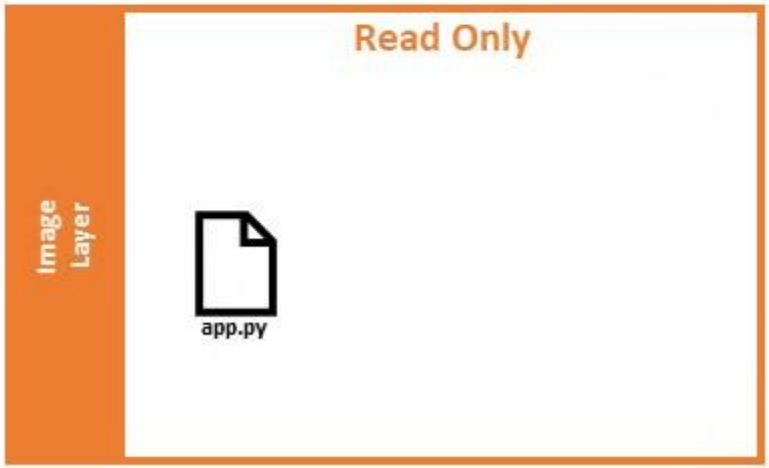

SkillAssure

# COPY-ON-WRITE

- Before I save the modified file, docker automatically creates a copy of the file in the read write layer
- This allow us to modify the application code
- All future modifications will be done on this copy of the file in the read write layer.
- This is called **copy on write mechanism**
- Image layer being read only just means that the files in these layers will not be modified in the image itself
- Image will remain the same all time until you rebuild

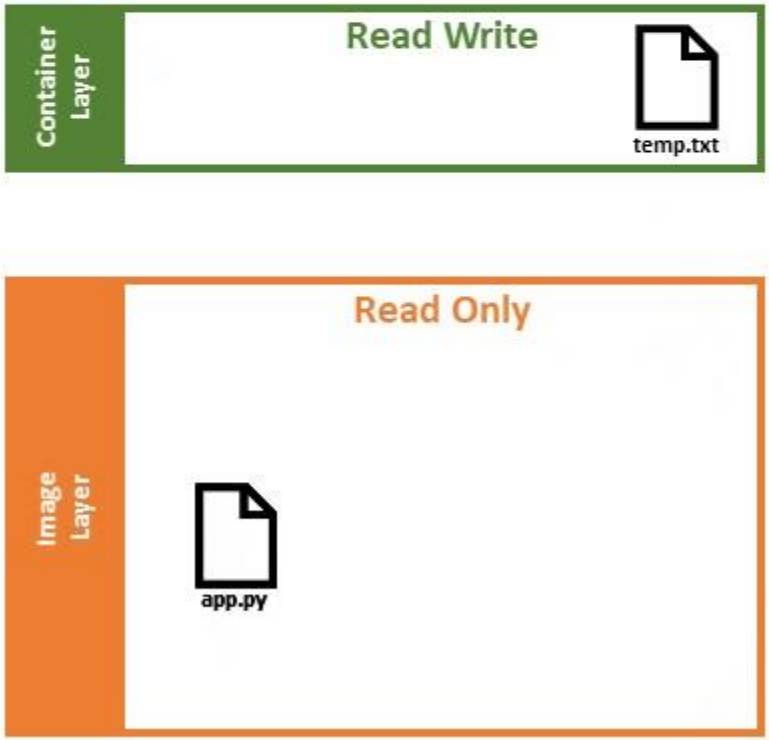

SkillAssure

# COPY-ON-WRITE

- What happens when I delete the container?
- All data stored in the container layer also gets deleted.
- The change we made to the app.py and temp.txt will also get removed.

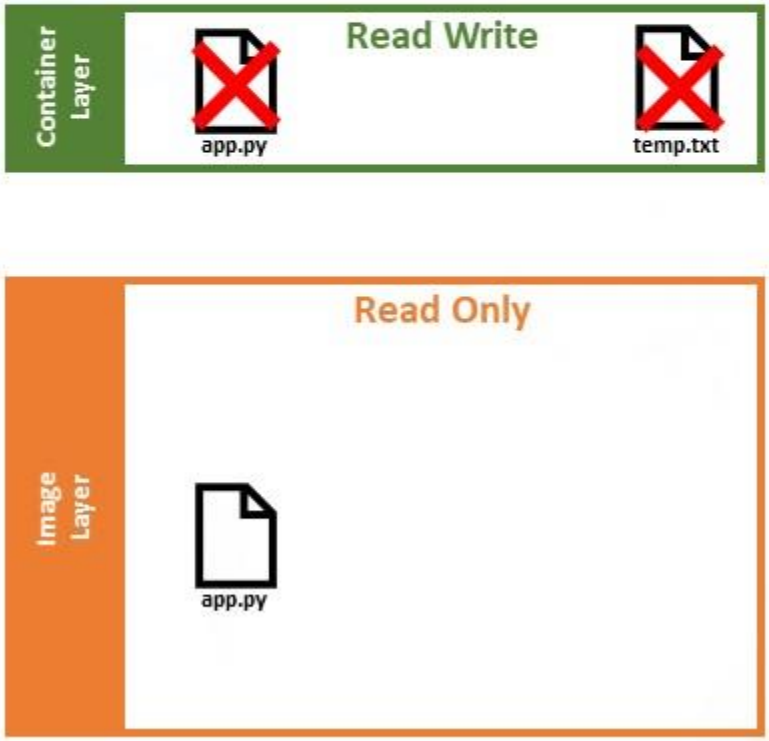

SkillAssure

# Volumes

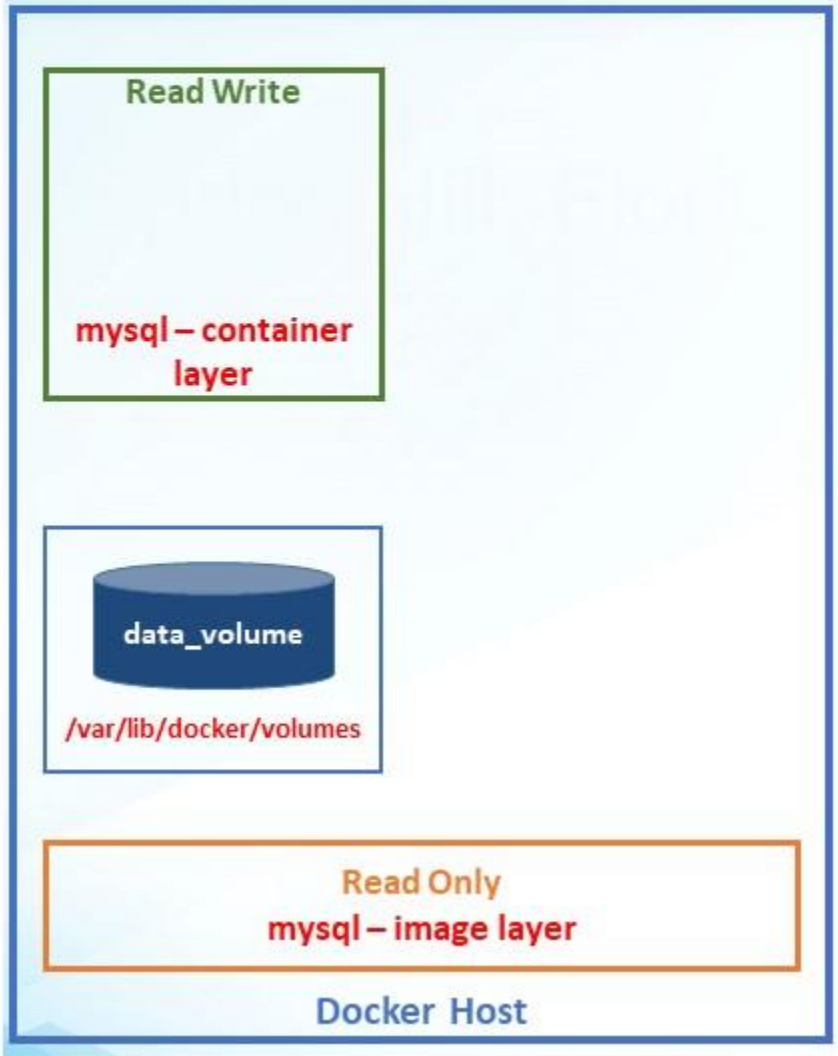


- How do we persist the data?
- Let's say I want to run MySQL and my databases to be persistent.

- We can add a persistent volume to the container
- Create a volume with name "data_volume"

```
docker volume create data_volume
```

- It creates a folder called data_volume under var/lib/docker/volumes

SkillAssure

# Volumes

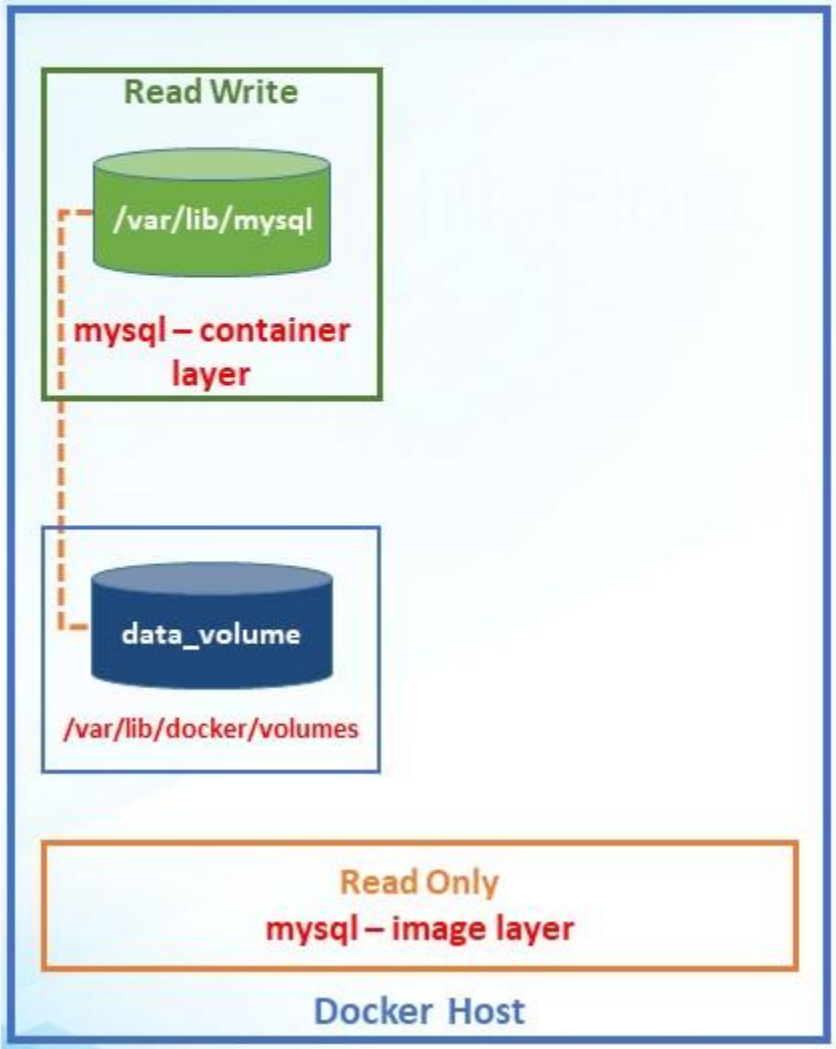


- **var/lib/mysql** is the default location where MySQL stores data
- To persist data I can mount this volume to MySql default location while running the container.

```
docker run -v data_volume:/var/lib/mysql mysql
```

- All data written on database is stored on the volume created on the docker host.
- Even if the container is destroyed the data will be available.

SkillAssure

# Volumes

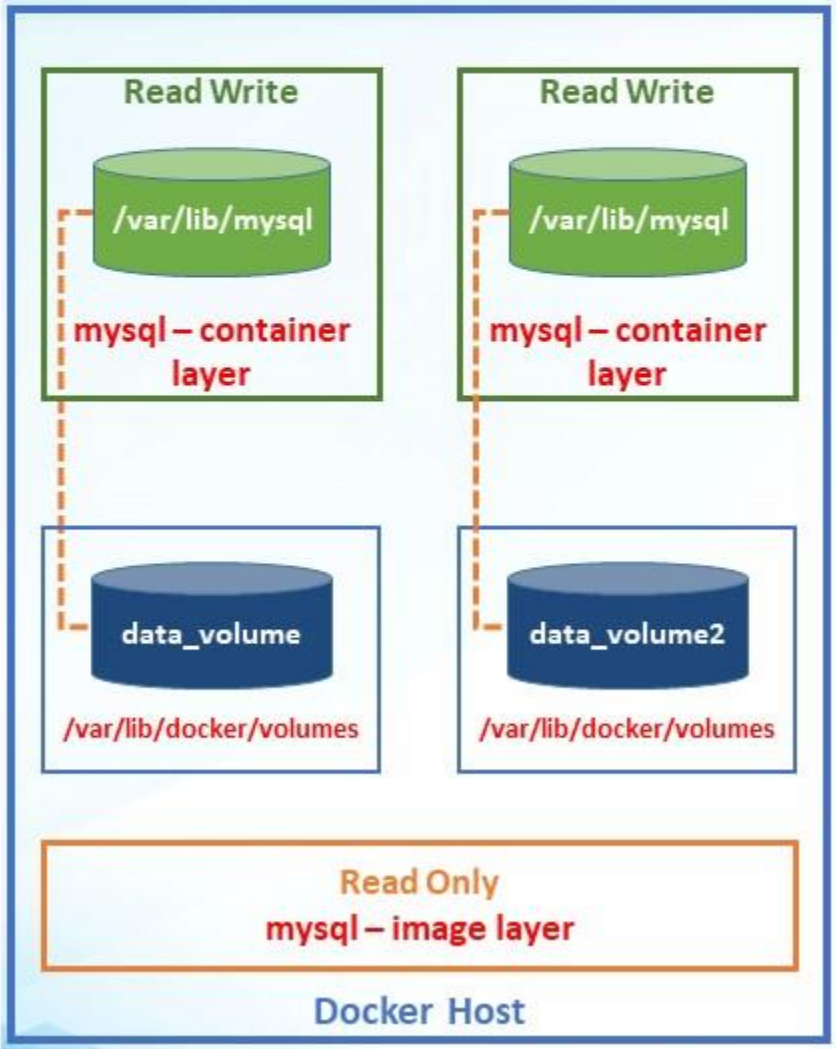


- What if you directly run the container without creating the volume?
- To persist data I can mount this volume to MySql default location while running the container.

```
docker run -v data_volume2:/var/lib/mysql mysql
```

- Docker will automatically create a volume named data_volume2 and mount it to the container.
- This is called volume mounting

SkillAssure

# Volumes

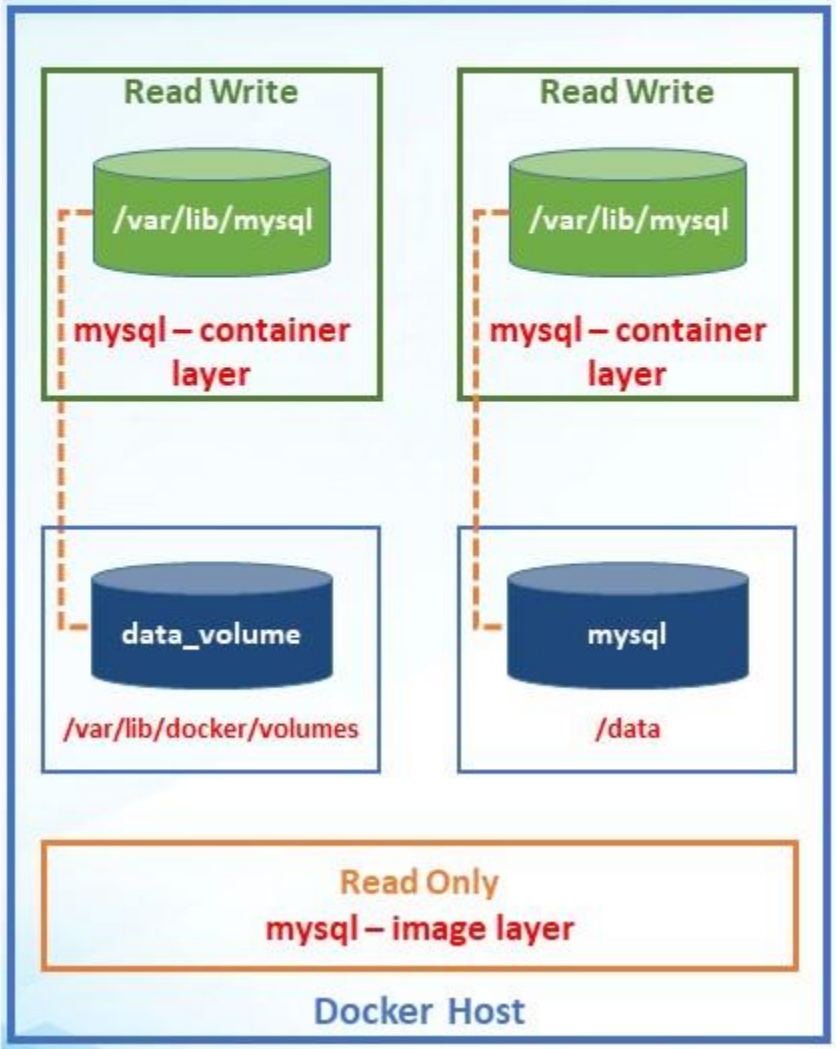


- What if our data is already present at another location, e.g. /data
- We would like to store database data on /data and not in the default /var/lib/docker/volumes folder.
- Create a container and provide complete path of folder to mount.

```
docker run -v /data/mysql:/var/lib/mysql mysql
```

- This is called **bind mounting**

SkillAssure

# Volumes

- –v is an old style for mounting

```
docker run -v /data/mysql:/var/lib/mysql mysql
```

- --mount is the preferred way as it is more verbose.
- Each parameter can be specified in a key=value format

```
docker run --mount type=bind,source=/data/mysql,target=/var/lib/mysql mysql
```

SkillAssure

# Volume Commands

- Create a volume

```
docker volume create my-vol
```

- List volumes

```
docker volume ls
```

- Inspect a volume:

```
docker volume inspect my-vol
```

- Remove a volume

```
docker volume rm my-vol
```

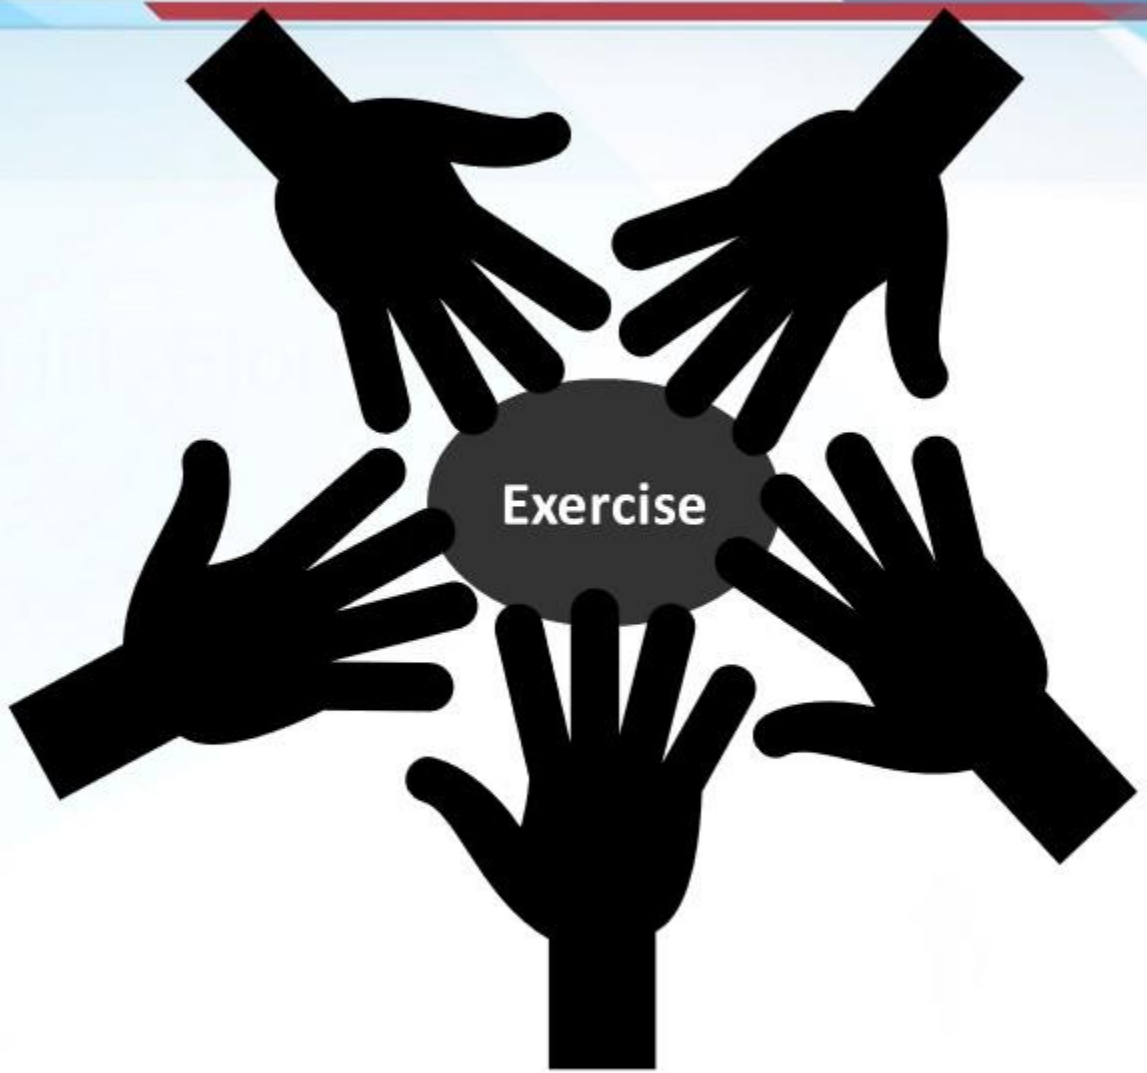
Exercise